

CP431 - Assignment 1: MPI Programming for Gaps Between Prime Numbers

Wilfrid Laurier Univeristy

Due Tuesday January 30, 2024

Torin Borton-McCallum - 190824620

Grant Westerholm - 190462000

Brandon Parker - 191593730

Rhea Sharma - 200576620

Taha Amir - 190728860

William Mabia - 190133240

Riley Adams - 190416070

Contents

1	Introduction	2
2	Algorithm Overview	2
2.1	Design Approach	2
2.2	Code Implementation	2
3	Algorithm Results	5
3.1	One Billion	5
3.2	One Trillion	5
4	Benchmarking	5
4.1	Scalability Analysis - One Billion	6
4.1.1	8 Processes	6
4.1.2	Increasing Processes (12, 16, 20)	6
4.1.3	Optimal Scaling (24 Processes)	6
4.1.4	Continued Improvement (28, 32 Processes)	6
4.2	MPI Wtime Resolution	6
4.3	One Trillion	6
5	Conclusion	7
6	Appendices	7
6.1	Run Times for Parallel Implementation	7
6.2	Run Time for Serial Implementation	8
6.3	Run Time for 1 trillion on Personal Computer	8

1 Introduction

This document presents an MPI (Message Passing Interface) program developed in C to determine the largest gap between a pair of consecutive prime numbers up to the ranges of one billion and one trillion. The objective is to explore parallel programming techniques for prime number computations on the SHARCnet cluster. The program utilizes Number Theoretic Functions from the GMP (GNU Multiple Precision Arithmetic Library) for accurate prime number calculations, specifically `mpz_probab_prime_p` and `mpz_nextprime`.

The MPI program is designed to execute on 2, 3, 4, 5, 6, 7, and 8 processors, and the performance is benchmarked to analyze its scalability. Results include the identification of the largest prime gap and the two prime numbers within the specified ranges that realize this gap. The documentation addresses code correctness, documentation quality, legibility, and overall presentation.

2 Algorithm Overview

2.1 Design Approach

1. **Range Calculation:** The global prime number range is initialized, and each MPI process is assigned a specific subrange for computation. The calculation is distributed to maximize parallelization, with each process focusing on its assigned segment.
2. **Prime Gap Calculation:**
The algorithm iterates through each subrange, identifying consecutive prime numbers and calculating the gaps between them. The largest gap, along with the primes before and after it, is tracked for each process.
3. **MPI Communication:**
To consolidate results, `MPI_Gather` is employed to collect information about the largest gaps and associated primes from each process. Process 0 then determines the overall largest gap and its corresponding prime numbers.
4. **Result Presentation:**
The final results, including the largest gap, the primes before and after the gap, the total number of processes involved, and computation times, are then formatted and displayed. GMP (GNU Multiple Precision) library functions facilitate the handling of large integers for precise calculations.

2.2 Code Implementation

```
1 // include statements
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <string.h>
5 #include <mpi.h>
6 #include <gmp.h>
7
8 // Using max fixed size instead of dynamic allocation because I kept running into
   buffer problems
9 #define MAX_STRING_LENGTH 1024
10
11 int main(int argc, char** argv) {
12
13     // Global variables to be used by each process
14     int world_rank, world_size;
15
16     MPI_Init(&argc, &argv);
17     MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
18     MPI_Comm_size(MPI_COMM_WORLD, &world_size);
19
20     // Range for each process
21     mpz_t range_start, range_end, global_range;
22
```

```

23 // Initialize and set the global range to 1 billion
24 mpz_init_set_ui(global_range, 1000000000);
25
26 mpz_t current_prime, next_prime, gap, max_gap, last_prime_in_range;
27 mpz_t prime_before_gap, prime_after_gap;
28
29 // Initialize all variables
30 mpz_init(current_prime);
31 mpz_init(next_prime);
32 mpz_init(gap);
33 mpz_init(max_gap);
34 mpz_init(last_prime_in_range);
35 mpz_init(prime_before_gap);
36 mpz_init(prime_after_gap);
37 mpz_set_ui(max_gap, 0);
38
39 // *****
40
41 // Calculate the range per process
42 unsigned long long int range_per_process = mpz_get_ui(global_range) / world_size;
43
44 // Set the start and end range for the current process
45 mpz_init_set_ui(range_start, world_rank * range_per_process + 1);
46 mpz_init_set_ui(range_end, (world_rank + 1) * range_per_process);
47
48 // Adjusted for last process - let's it cover entire range
49 if (world_rank == world_size - 1) {
50     mpz_set(range_end, global_range);
51 }
52
53 // Synchronize all processes
54 MPI_Barrier(MPI_COMM_WORLD);
55
56 // Get the starting time
57 double start_time = MPI_Wtime();
58
59 // Get the next prime starting from range_start
60 mpz_nextprime(current_prime, range_start);
61
62 while (mpz_cmp(current_prime, range_end) <= 0) {
63
64     // Get next prime and calculate gap
65     mpz_nextprime(next_prime, current_prime);
66
67     // Calculate the gap between the current and next prime
68     mpz_sub(gap, next_prime, current_prime);
69
70     // store if largest gap
71     if (mpz_cmp(gap, max_gap) > 0) {
72         mpz_set(max_gap, gap);
73         mpz_set(prime_before_gap, current_prime);
74         mpz_set(prime_after_gap, next_prime);
75     }
76
77     // Move to the next prime
78     mpz_set(current_prime, next_prime);
79 }
80
81 // Set the last prime in the range
82 mpz_set(last_prime_in_range, current_prime);
83
84 // Synchronize all processes
85 MPI_Barrier(MPI_COMM_WORLD);
86
87 // Get the ending time
88 double end_time = MPI_Wtime();
89
90 // Calculate the elapsed time
91 double elapsed_time = end_time - start_time;
92
93 // Get the resolution of MPI_Wtime

```

```

94     double tick = MPI_Wtick();
95
96     // *****
97
98     // Prepare strings for communication
99     char max_gap_str[MAX_STRING_LENGTH], prime_before_str[MAX_STRING_LENGTH],
100     prime_after_str[MAX_STRING_LENGTH];
101
102     // Initialize strings
103     memset(max_gap_str, 0, MAX_STRING_LENGTH);
104     memset(prime_before_str, 0, MAX_STRING_LENGTH);
105     memset(prime_after_str, 0, MAX_STRING_LENGTH);
106     mpz_get_str(max_gap_str, 10, max_gap);
107
108     // Convert prime_before_gap and prime_after_gap to string
109     mpz_get_str(prime_before_str, 10, prime_before_gap);
110     mpz_get_str(prime_after_str, 10, prime_after_gap);
111
112     // Allocate memory for string receive in process 0
113     char (*all_max_gap_str)[MAX_STRING_LENGTH] = NULL;
114     char (*all_prime_before_str)[MAX_STRING_LENGTH] = NULL;
115     char (*all_prime_after_str)[MAX_STRING_LENGTH] = NULL;
116
117     if (world_rank == 0) {
118         // all_max_gap_str = malloc(world_size * MAX_STRING_LENGTH * sizeof(char));
119         // all_prime_before_str = malloc(world_size * MAX_STRING_LENGTH * sizeof(char))
120     );
121         // all_prime_after_str = malloc(world_size * MAX_STRING_LENGTH * sizeof(char))
122         ;
123
124         // This lets it compile with mpi c++ above only compiles with mpic
125         all_max_gap_str = (char (*)[MAX_STRING_LENGTH])malloc(world_size *
126         MAX_STRING_LENGTH * sizeof(char));
127         all_prime_before_str = (char (*)[MAX_STRING_LENGTH])malloc(world_size *
128         MAX_STRING_LENGTH * sizeof(char));
129         all_prime_after_str = (char (*)[MAX_STRING_LENGTH])malloc(world_size *
130         MAX_STRING_LENGTH * sizeof(char));
131     }
132
133     // Gather fixed size strings
134     MPI_Gather(max_gap_str, MAX_STRING_LENGTH, MPI_CHAR, all_max_gap_str,
135     MAX_STRING_LENGTH, MPI_CHAR, 0, MPI_COMM_WORLD);
136     MPI_Gather(prime_before_str, MAX_STRING_LENGTH, MPI_CHAR, all_prime_before_str,
137     MAX_STRING_LENGTH, MPI_CHAR, 0, MPI_COMM_WORLD);
138     MPI_Gather(prime_after_str, MAX_STRING_LENGTH, MPI_CHAR, all_prime_after_str,
139     MAX_STRING_LENGTH, MPI_CHAR, 0, MPI_COMM_WORLD);
140
141     if (world_rank == 0) {
142         mpz_t global_max_gap, global_prime_before_gap, global_prime_after_gap;
143
144         // Initialize global values
145         mpz_init(global_max_gap);
146         mpz_init(global_prime_before_gap);
147         mpz_init(global_prime_after_gap);
148         mpz_set_ui(global_max_gap, 0);
149
150         // Process received strings
151         for (int i = 0; i < world_size; ++i) {
152             mpz_t temp_max_gap;
153             mpz_init(temp_max_gap);
154             mpz_set_str(temp_max_gap, all_max_gap_str[i], 10);
155
156             // gmp_printf("temp_max_gap: %Zd\n", temp_max_gap);
157             // gmp_printf("global_max_gap: %Zd\n", global_max_gap);
158
159             if (mpz_cmp(temp_max_gap, global_max_gap) > 0) {
160                 mpz_set(global_max_gap, temp_max_gap);
161                 mpz_set_str(global_prime_before_gap, all_prime_before_str[i], 10);
162                 mpz_set_str(global_prime_after_gap, all_prime_after_str[i], 10);
163             }
164         }
165     }

```

```

156         // Clear temporary maximum gap
157         mpz_clear(temp_max_gap);
158     }
159
160
161     // gmp_printf("Largest gap: %Zd, between primes %Zd and %Zd\n", max_gap,
prime_before_gap, prime_after_gap);
162     gmp_printf("Largest gap: %Zd, between primes %Zd and %Zd\n", global_max_gap,
global_prime_before_gap, global_prime_after_gap);
163     printf("Total processes: %d\n", world_size);
164     printf("Total computation time: %e seconds\n", elapsed_time);
165     printf("Average computation time per process: %e seconds\n", elapsed_time /
world_size);
166     printf("Resolution of MPI_Wtime: %e seconds\n", tick);
167
168     // Cleanup
169     mpz_clear(global_max_gap);
170     mpz_clear(global_prime_before_gap);
171     mpz_clear(global_prime_after_gap);
172     free(all_max_gap_str);
173     free(all_prime_before_str);
174     free(all_prime_after_str);
175 }
176
177 // Cleanup
178 mpz_clear(range_start);
179 mpz_clear(range_end);
180 mpz_clear(global_range);
181 mpz_clear(current_prime);
182 mpz_clear(next_prime);
183 mpz_clear(gap);
184 mpz_clear(max_gap);
185 mpz_clear(last_prime_in_range);
186 mpz_clear(prime_before_gap);
187 mpz_clear(prime_after_gap);
188
189 MPI_Finalize();
190 return 0;
191 }

```

3 Algorithm Results

3.1 One Billion

Running the algorithm listed above, the largest gap we produced between two contiguous prime numbers is 282. The two prime numbers that materialize this gap are 436273009 and 436273291. This is consistent with the result produced from a relatively equivalent serially programmed algorithm to accomplish the same task. The serial version took approximately 7 times longer to run than the parallel implementation with 8 processes.

3.2 One Trillion


From one billion to one trillion, we expanded the algorithm to look for a larger gap between two contiguous prime numbers. The largest gap found increased from 282 to 540. The two prime numbers that materialize this largest gap are 738832927927 and 738832928467. Trying to run a serial implementation on one trillion is not feasible since the parallel solution's runtime was around 5 hours.

4 Benchmarking

The benchmarking results presented here evaluate the performance of a parallelized algorithm for computing prime gaps. The primary metric considered is the largest observed gap, which measures 282 between primes 436273009 and 436273291. The computations were conducted using MPI, a widely used message-passing interface for parallel computing.

4.1 Scalability Analysis - One Billion

Parallel Execution Time Scaling with Processors for Largest Gap Between One Billion



4.1.1 8 Processes

In the initial run with 8 processes, the total computation time was 124.63 seconds, yielding an average computation time per process of 15.58 seconds. This serves as a baseline for evaluating the scalability of the algorithm.

4.1.2 Increasing Processes (12, 16, 20)

As the number of processes increased to 12, 16, and 20, there was a notable improvement in performance. The total computation time decreased to 87.67 seconds, 69.32 seconds, and 56.39 seconds, respectively, showcasing the scalability of the parallelized algorithm.

4.1.3 Optimal Scaling (24 Processes)

Further scaling the computation to 24 processes increased efficiency, resulting in a total computation time of 47.44 seconds and an average computation time per process of 1.98 seconds. This suggests a near-linear scaling, indicative of effective parallelization.

4.1.4 Continued Improvement (28, 32 Processes)

Continuing the trend, 28 processes led to a total computation time of 41.68 seconds, and with 32 processes, the total computation time reduced to 36.69 seconds. The average computation time per process dropped to 1.49 seconds and 1.15 seconds, respectively, emphasizing the continued improvement with an increased number of processes.

4.2 MPI Wtime Resolution

It is important to note that the resolution of MPI Wtime, a function measuring time in MPI programs, remained consistent at 1.000000e-09 seconds across all experiments. This consistent resolution highlights the accuracy of the timing measurements.

4.3 One Trillion

This computation involved 10 processes, resulting in a total computation time of 16,651.70 seconds, translating to an average computation time per process of 1,665.17 seconds. The resolution of MPI Wtime for this experiment was set at 1.000000e-06 seconds. Comparing these results with the previous benchmarking conducted within the billion range, where the largest gap was 282, there is a noticeable

increase in both the size of the gap and the computation time. The larger search space in the trillion range naturally leads to more extensive computations, reflected in the increased gap size and the higher overall computation time. This highlights the scalability of the algorithm to handle significantly larger ranges.

5 Conclusion

In conclusion, the code implemented an MPI program in C designed to find the largest gap between consecutive prime numbers within the ranges of one billion and one trillion. The use of the GMP library for precise calculations such as the range and prime gap calculations, allows us to explore parallel programming by utilizing two to eight processors and assessing scalability through benchmarking. The algorithm efficiently distributes computations, identifies prime numbers and tracks gaps, consolidating results with `MPI_Gather`. The results demonstrate efficient parallelization, with a largest gap of 282 in the billion range and 540 in the trillion range. The scalability in terms of this algorithm is evident, with reduced computation time as the number of processes increases. The consistent resolution of `MPI_Wtime` emphasizes accurate timing measurements and highlights the algorithm's reliability across varying computational scales. Overall, this research showcases the algorithm's scalability, efficiently handling larger search spaces and substantiating its effectiveness in parallelized prime number computations.

6 Appendices

6.1 Run Times for Parallel Implementation

Largest gap: 282, between primes 436273009 and 436273291
Total processes: 8
Total computation time: 1.246293e+02 seconds
Average computation time per process: 1.557866e+01 seconds
Resolution of `MPI_Wtime`: 1.000000e-09 seconds

Largest gap: 282, between primes 436273009 and 436273291
Total processes: 12
Total computation time: 8.767491e+01 seconds
Average computation time per process: 7.306242e+00 seconds
Resolution of `MPI_Wtime`: 1.000000e-09 seconds

Largest gap: 282, between primes 436273009 and 436273291
Total processes: 16
Total computation time: 6.932332e+01 seconds
Average computation time per process: 4.332707e+00 seconds
Resolution of `MPI_Wtime`: 1.000000e-09 seconds

Largest gap: 282, between primes 436273009 and 436273291
Total processes: 20
Total computation time: 5.639275e+01 seconds
Average computation time per process: 2.819637e+00 seconds
Resolution of `MPI_Wtime`: 1.000000e-09 seconds

Largest gap: 282, between primes 436273009 and 436273291
Total processes: 24
Total computation time: 4.744021e+01 seconds
Average computation time per process: 1.976675e+00 seconds
Resolution of `MPI_Wtime`: 1.000000e-09 seconds

Largest gap: 282, between primes 436273009 and 436273291
Total processes: 28

Total computation time: 4.167845e+01 seconds
Average computation time per process: 1.488516e+00 seconds
Resolution of MP_Wtime: 1.000000e-09 seconds

Largest gap: 282, between primes 436273009 and 436273291
Total processes: 32
Total computation time: 3.668536e+01 seconds
Average computation time per process: 1.146418e+00 seconds
Resolution of MPI_Wtime: 1.000000e-09 seconds

6.2 Run Time for Serial Implementation

Time taken: 938.300000 seconds
lprime: 436273009 rprime: 43627329821 gap: 282

6.3 Run Time for 1 trillion on Personal Computer

Largest gap: 540, between primes 738832927927 and 738832928467
Total processes: 10
Total computation time: 1.665170e+04 seconds
Average computation time per process: 1.665170e+03 seconds
Resolution of MPI_Wtime: 1.000000e-06 seconds